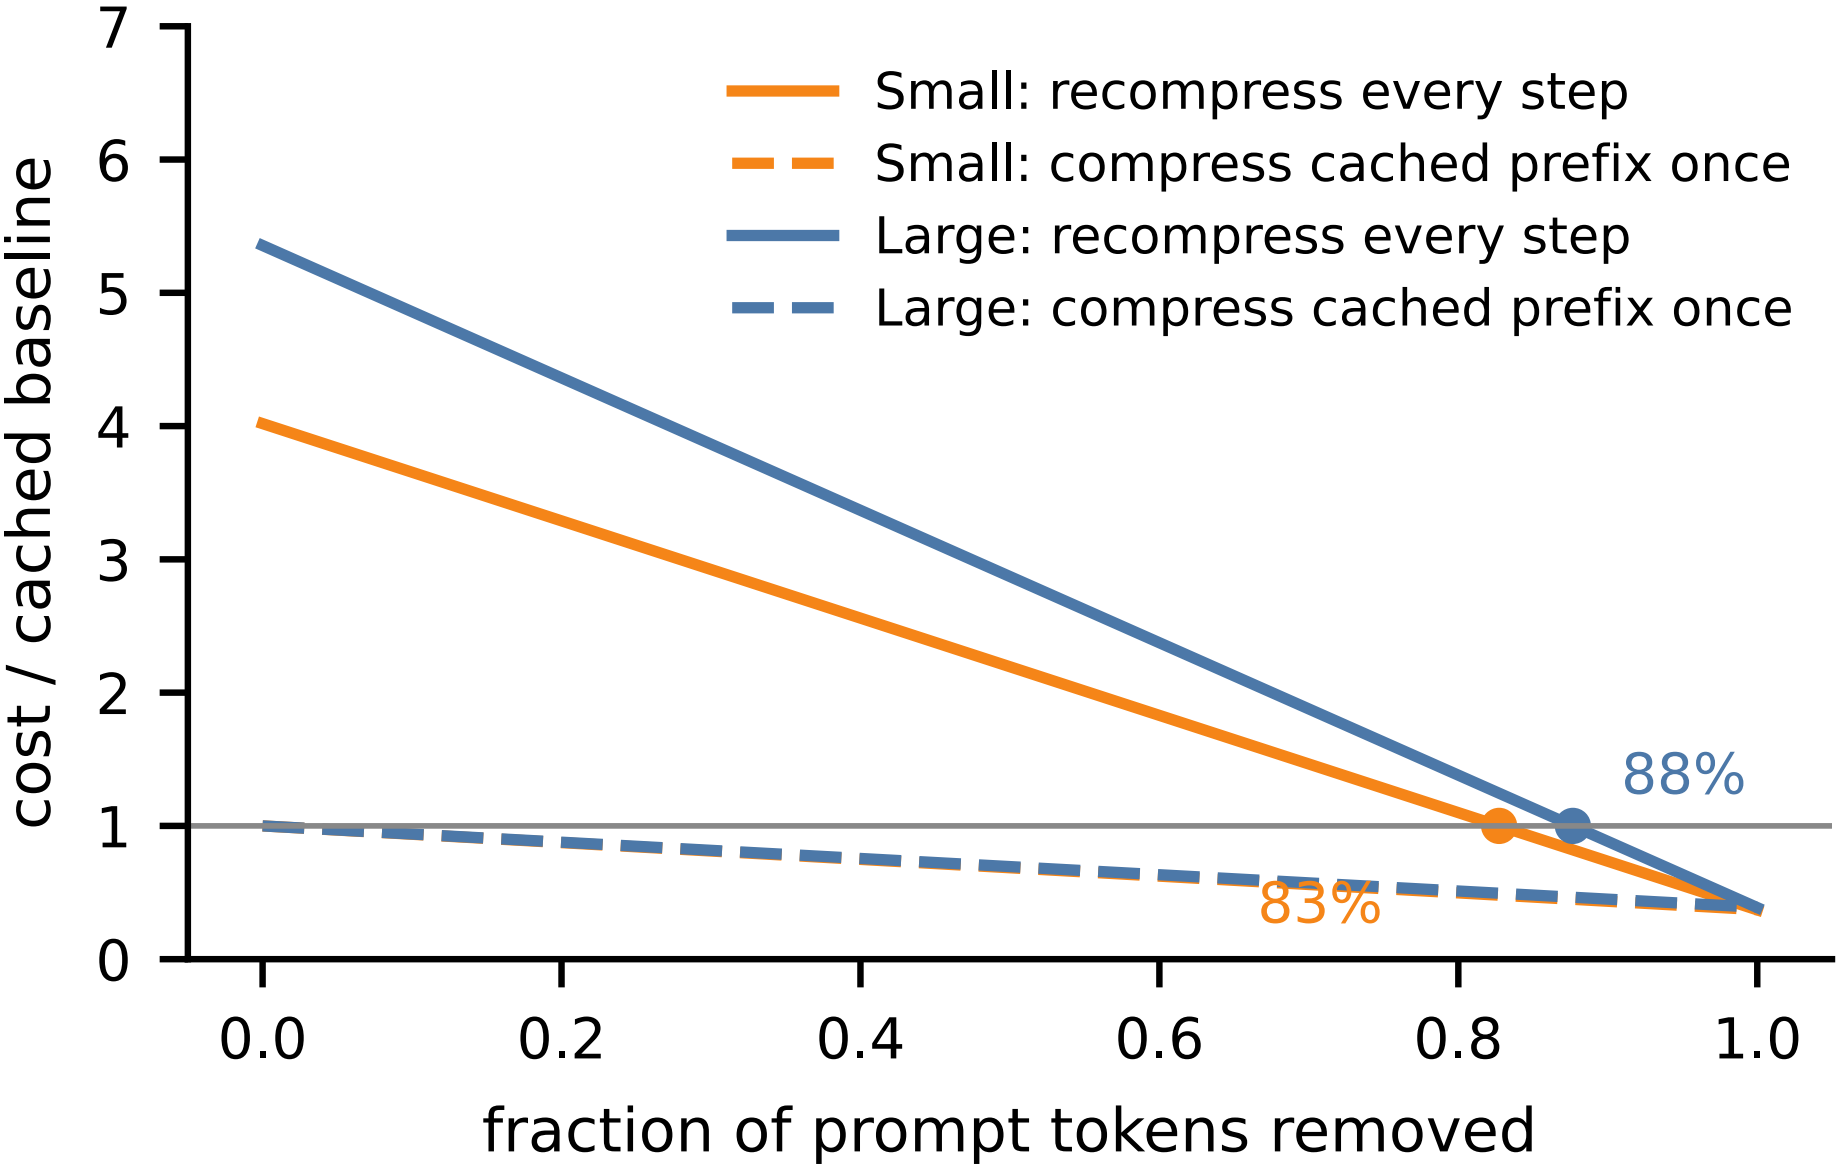
Small: recompress every step
Small: compress cached prefix once
Large: recompress every step
Large: compress cached prefix once
cost / cached baseline
fraction of prompt tokens removed
0
1
2
3
4
5
6
7
0.0
0.2
0.4
0.6
0.8
1.0
88%
83%